

http://

PART 1: GOING DEEP ON HTTP  
PART 2: HIGH LEVEL HTTP APIs

Jesse Wilson



# Hi, I'm Jesse

---

- Infrastructure nerd
- Former Android core libraries lead
  - Maintained Android's HttpURLConnection and Apache HTTP client
- Today I'm at Square
  - Working on OkHttp & MockWebServer

# Plan

---

- HTTP in 2013
- HTTP clients for Android
  - Apache HTTP client
  - HttpURLConnection
  - OkHttp
- Essential dev tools
  - Charles Web Debugging Proxy
  - MockWebServer

# Plan

---

- Retrofit
- Picasso
- Volley

# Follow along

<http://tinyurl.com/androidhttp>

The screenshot shows a web browser window displaying a GitHub repository page. The URL in the address bar is <https://github.com/swankjesse/android-http-examples/tree/master/src/main/java/com/publicobject/http>. The page title is "swankjesse / android-http-examples". The repository is public. The "Code" tab is selected. The "branch: master" dropdown shows "master". Below the tabs, there are links for "Files", "Commits", "Branches", and "Tags". The main content area shows a commit history for the "http" directory. The first commit is by "square-build-bot" and was authored 6 minutes ago. The commit message is "Initial API examples". The latest commit is also by "square-build-bot" and has the commit ID "fb97ee1485". Below the commit history, there are two more commits: "cache" and "get", both authored 6 minutes ago by "square-build-bot" with the message "Initial API examples [square-build-bot]".

Commit	Author	Date	Message
Initial API examples	square-build-bot	6 minutes ago	Initial API examples [square-build-bot]
cache	square-build-bot	6 minutes ago	Initial API examples [square-build-bot]
get	square-build-bot	6 minutes ago	Initial API examples [square-build-bot]

HTTP in 2013

# Demo

# Port 80

---

```
$ telnet swank.ca 80
Trying 69.163.193.18...
Connected to swank.ca.
Escape character is '^]'.
GET / HTTP/1.1
Host: swank.ca

HTTP/1.1 200 OK
Date: Tue, 08 Jan 2013 00:16:56 GMT
Server: Apache
Last-Modified: Wed, 14 Sep 2011 19:12:02 GMT
ETag: "27b-4aceb88f6a480"
Accept-Ranges: bytes
Content-Length: 199
Vary: Accept-Encoding
Content-Type: text/html; charset=utf-8

<html>
  <head><title>swank.ca</title></head>
  <body align="center" bgcolor="#003399">
    
  </body>
</html>
```

# Proxies

---

- HTTP proxies exist, annoy
- You speak HTTP to the proxy. It speaks HTTP to the world.
- HTTP and HTTPS

# Caching

---

- Fundamental
- HTTP distinguishes between resources (GET) and actions (POST).

# Caching example

---

# Caching: for servers

---

- Use the Cache-Control & ETag headers.

HTTP/1.1 200 OK

Cache-Control: max-age=3600

Etag: v3

...

- Using modification dates is clumsy! The user agent is dumb and uses heuristics.

# Caching: for servers

---

- Serving static resources like images?  
Use a permanent URL and let it cache forever!
- When the resource is updated, just change the URL!

# Content Delivery Networks (CDNs)

- Occasionally, you're bottlenecked by the speed of light



# Thumbor

---



<http://im.squareup.com/stand.png>



<http://thumbor.squareup.com/unsafe/160x160/http://im.squareup.com/stand.png>

# Gzip

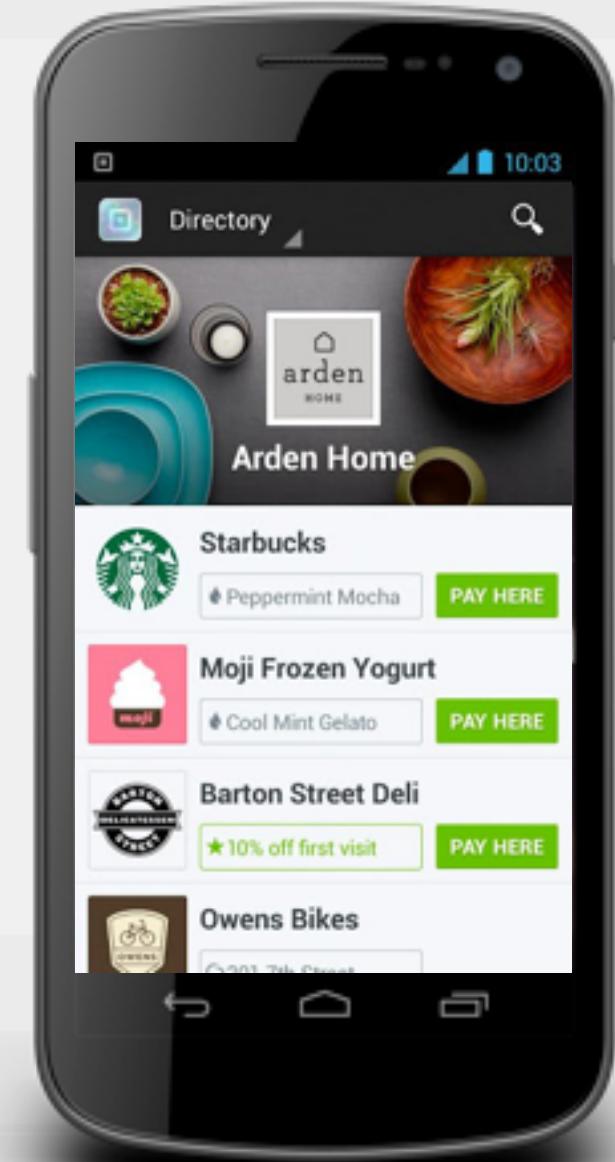
---

JSON: **53** KiB

Proto: **31** KiB

gzipped JSON: **9.1** KiB

gzipped Proto: **8.5** KiB



## Gzip: Clients

---

- Use the Accept-Encoding request header:

```
GET /directory.json HTTP/1.1  
Accept-Encoding: gzip
```

...

# Gzip: Servers

---

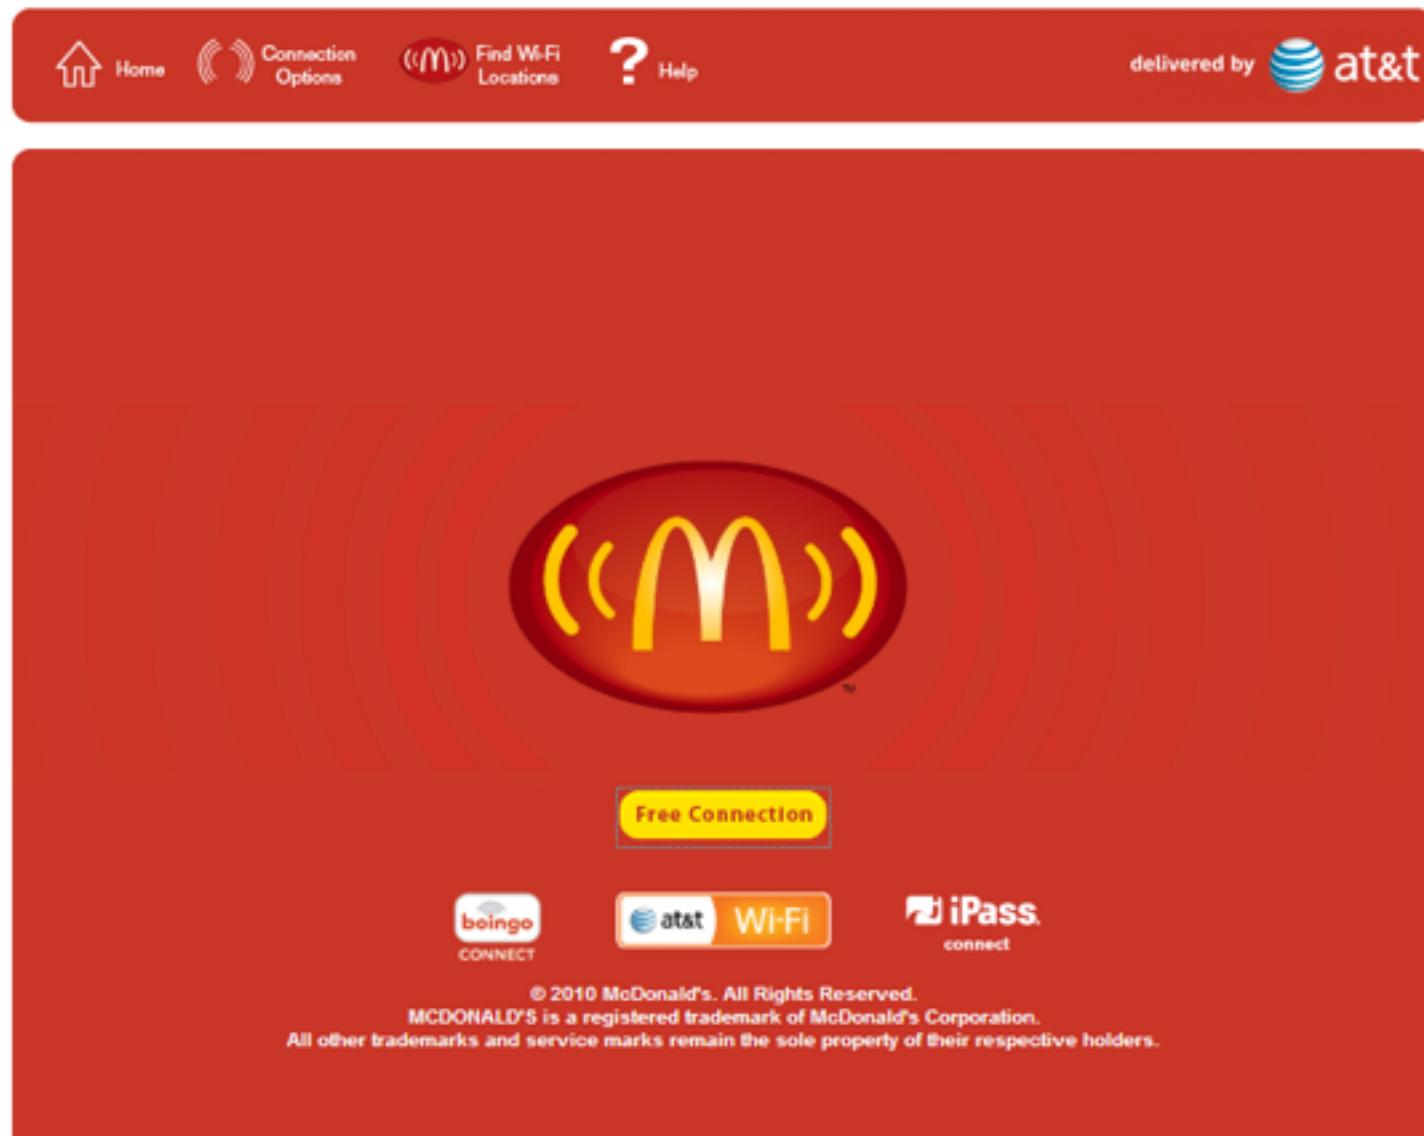
- Using gzip on already-compressed content like images is wasteful. Don't double compress content types like image/png or image/jpeg

# DNS

---

- Redundant datacenters?
- Supporting IPv6 and IPv4?
- Each host may have multiple IP addresses!

# WiFi Hotspots



The image shows the McDonald's WiFi Hotspots landing page. The background is red with the iconic golden arches logo in the center. At the top, there is a navigation bar with four links: "Home" (with a house icon), "Connection Options" (with a signal icon), "Find Wi-Fi Locations" (with a location pin icon), and "Help" (with a question mark icon). To the right of the navigation bar is the text "delivered by" followed by the AT&T logo. Below the navigation bar is a large yellow button with the text "Free Connection". At the bottom of the page, there are three logos: "boingo CONNECT", "at&t Wi-Fi", and "iPass connect". A small copyright notice at the very bottom reads: "© 2010 McDonald's. All Rights Reserved. MCDONALD'S is a registered trademark of McDonald's Corporation. All other trademarks and service marks remain the sole property of their respective holders."

Home Connection Options Find Wi-Fi Locations Help

delivered by  at&t

Free Connection

boingo CONNECT at&t Wi-Fi iPass connect

© 2010 McDonald's. All Rights Reserved.  
MCDONALD'S is a registered trademark of McDonald's Corporation.  
All other trademarks and service marks remain the sole property of their respective holders.

# WiFi Hotspots

---

- Coffee shop?
- Airport?
- Your HTTP requests may be redirected.

# HTTPS

---

- Transport Layer Security (TLS), the successor to Secure Sockets Layer (SSL).
- Reduces eavesdropping while you use the WiFi at Panera Bread or the airport

# Certificate Authorities

---

- Lots of problems
- Governments want to spy

# Certificate Authorities

---

- Browser vendors have two bad options:
  - Include a CA
  - Don't include a CA

# Certificate Authorities

---

- Hardcode your cert in your app:

```
SSLocketFactory goodSslSocketFactory(Context context) {  
    KeyStore trusted = KeyStore.getInstance("BKS");  
    InputStream in = context.getResources().openRawResource(R.raw.truststore);  
    trusted.load(in, TRUST_STORE_PASSWORD);  
    SSLContext sslContext = SSLContext.getInstance("TLS");  
    TrustManagerFactory trustManagerFactory = TrustManagerFactory.getInstance(  
        TrustManagerFactory.getDefaultAlgorithm());  
    trustManagerFactory.init(trusted);  
    sslContext.init(null, trustManagerFactory.getTrustManagers(), null);  
    return sslContext.getSocketFactory();  
}
```

# Hostname Verification

---

- Does \*.squareup.com match beta.api.squareup.com ?
- Validate that the certificate matches the host.
- Rules followed inconsistently!

# The Most Dangerous Code in the World: Validating SSL Certificates in Non-Browser Software

Martin Georgiev  
The University of Texas  
at Austin

Rishita Anubhai  
Stanford University

Subodh Iyengar  
Stanford University

Dan Boneh  
Stanford University

Suman Jana  
The University of Texas  
at Austin

Vitaly Shmatikov  
The University of Texas  
at Austin

## ABSTRACT

SSL (Secure Sockets Layer) is the de facto standard for secure Internet communications. Security of SSL connections against an active network attacker depends on correctly validating public-key certificates presented when the connection is established.

We demonstrate that SSL certificate validation is completely broken in many security-critical applications and libraries. Vulnerable software includes Amazon’s EC2 Java library and all cloud clients based on it; Amazon’s and PayPal’s merchant SDKs responsible for transmitting payment details from e-commerce sites to payment gateways; integrated shopping carts such as osCommerce, ZenCart, Ubercart, and PrestaShop; AdMob code used by mobile websites; Chase mobile banking and several other Android apps and libraries; Java Web-services middleware—including Apache Axis, Axis 2, Codehaus XFire, and Pusher library for Android—and *all* applications employing this middleware. Any SSL connection from any of

cations. The main purpose of SSL is to provide end-to-end security against an active, man-in-the-middle attacker. Even if the network is completely compromised—DNS is poisoned, access points and routers are controlled by the adversary, etc.—SSL is intended to guarantee confidentiality, authenticity, and integrity for communications between the client and the server.

Authenticating the server is a critical part of SSL connection establishment.<sup>1</sup> This authentication takes place during the SSL handshake, when the server presents its public-key certificate. In order for the SSL connection to be secure, the client must carefully verify that the certificate has been issued by a valid certificate authority, has not expired (or been revoked), the name(s) listed in the certificate match(es) the name of the domain that the client is connecting to, and perform several other checks [14, 15].

SSL implementations in Web browsers are constantly evolving through “penetrate-and-patch” testing, and many SSL-related vulnerabilities in browsers have been repaired over the years. SSL

# Cookies

---

- Four specs:  
Netscape, RFC 2109, RFC 2965, RFC 6265

*"Prior to this document, there were at least three descriptions of cookies."*

*"However, none of these documents describe how the Cookie and Set-Cookie headers are actually used on the Internet."*

# Cookies

---

- Awkward on Android. No out-of-the-box place to persist across app launches.

# Ranges

---

- Restarting a failed download sucks.
- If your app does large downloads (> 5 MiB) you should look into the Range header. You'll need to do this manually.

# Sockets

---

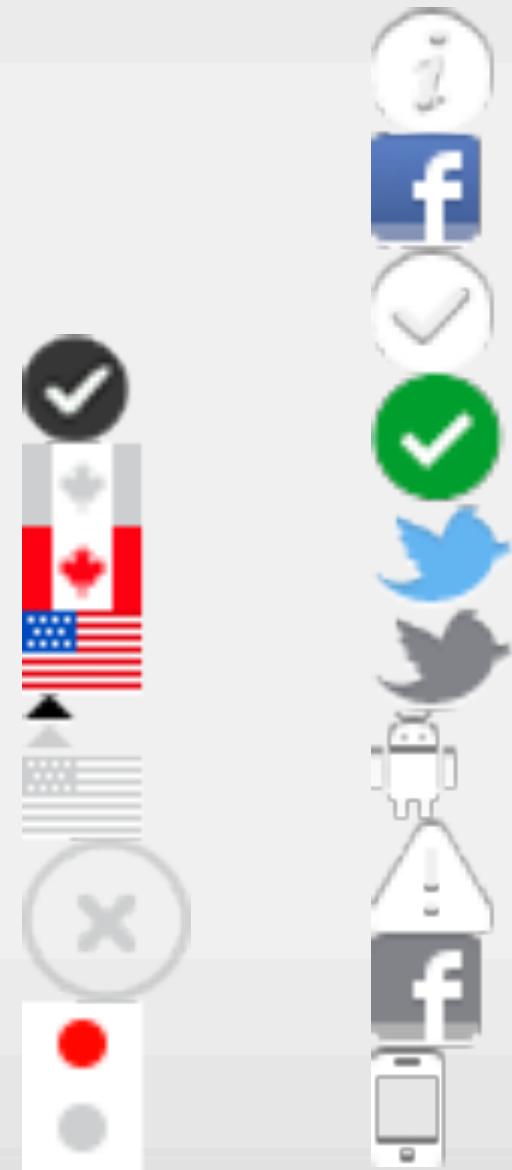
*"A single-user client SHOULD NOT maintain more than 2 connections with any server or proxy."*

- IE 10 uses 8
- Chrome uses 6
- Firefox uses 6

# Sockets

---

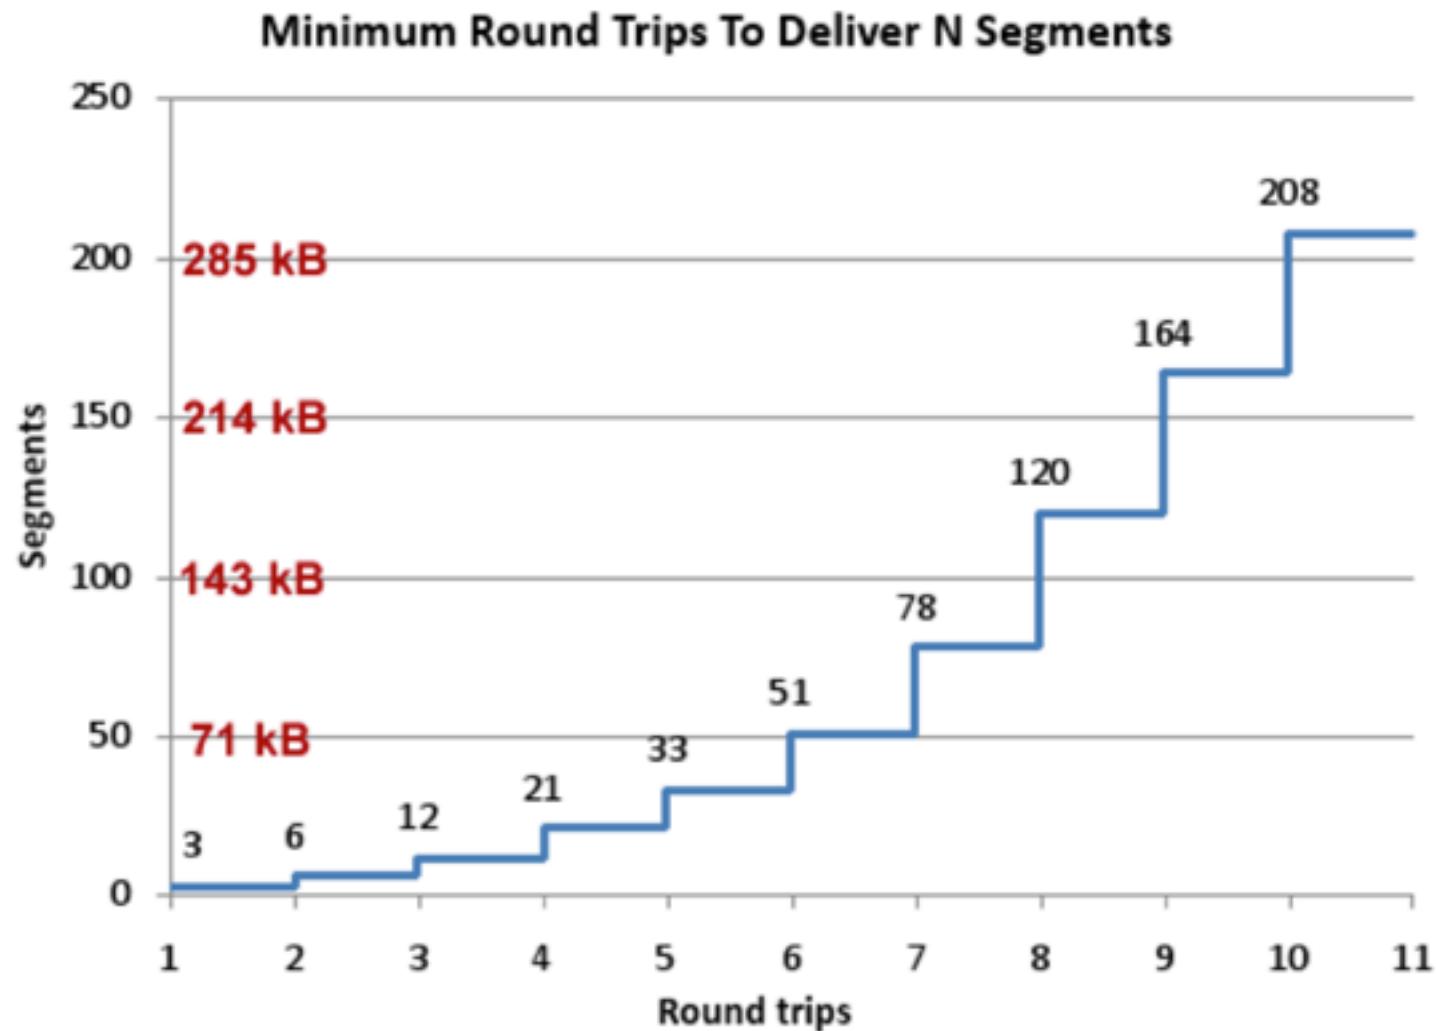
- Hacks!
  - Image spriting
  - Flattening CSS
  - Google serves content from multiple hostnames!
- These hacks hurt caching!



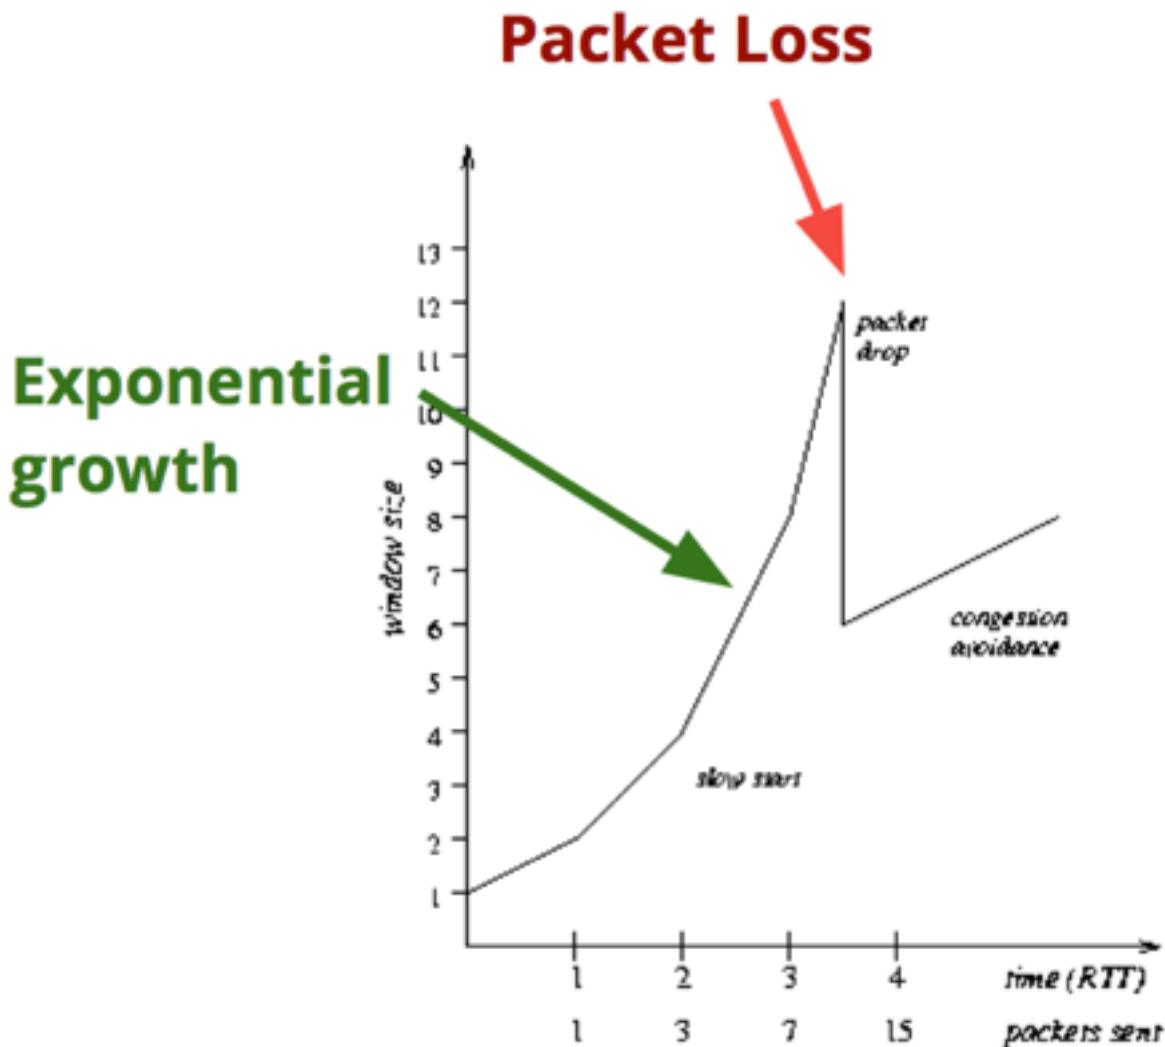
locale.png

icon.png

# Sockets



# Sockets



# Sockets

---

- Connection pooling

# SPDY

---

- Keep the good parts!
  - But fix the sockets problem
  - Compress more stuff
- 
- Reasonable go-to-market strategy:  
Chrome browser + google.com website + NPN

# SPDY

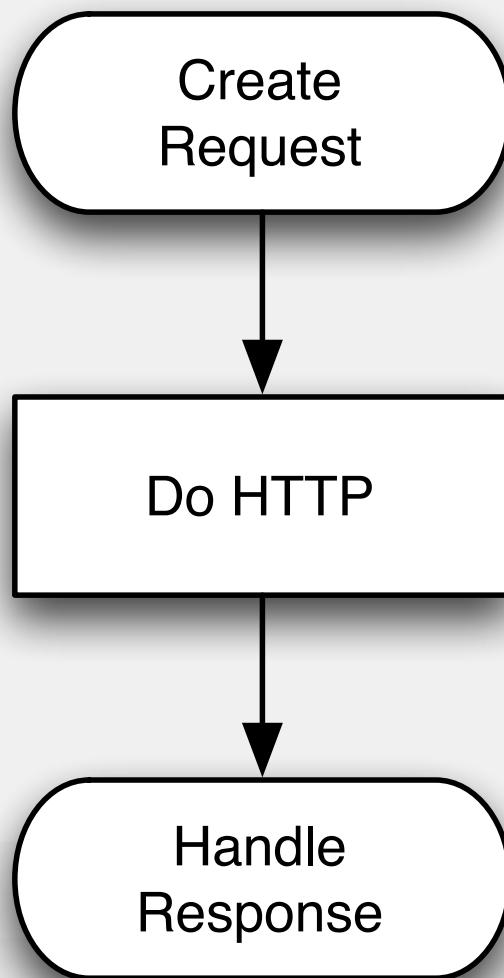
---

- Needs HTTPS
- Needs Android 4.1+

# Android's HTTP Clients

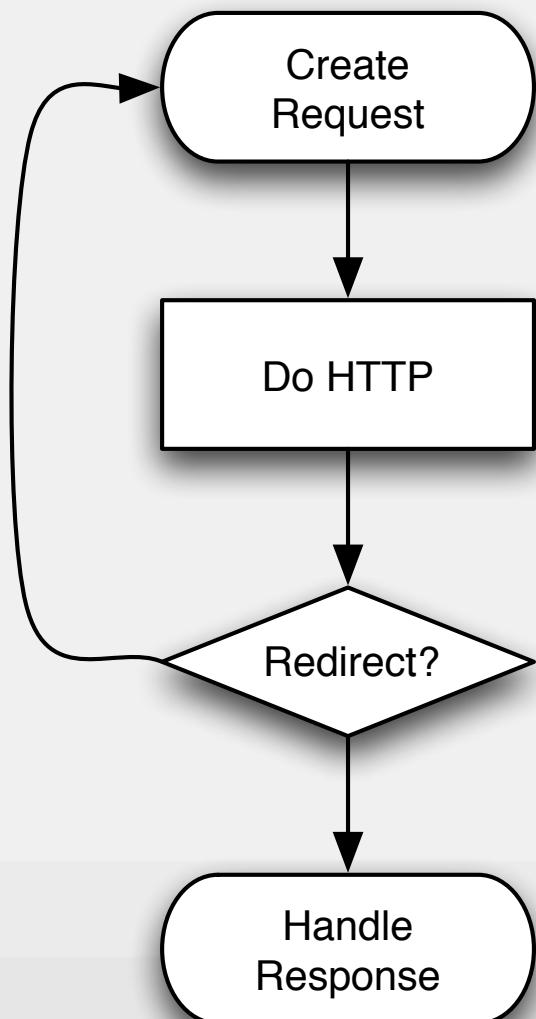
# HTTP

---



# HTTP with redirects

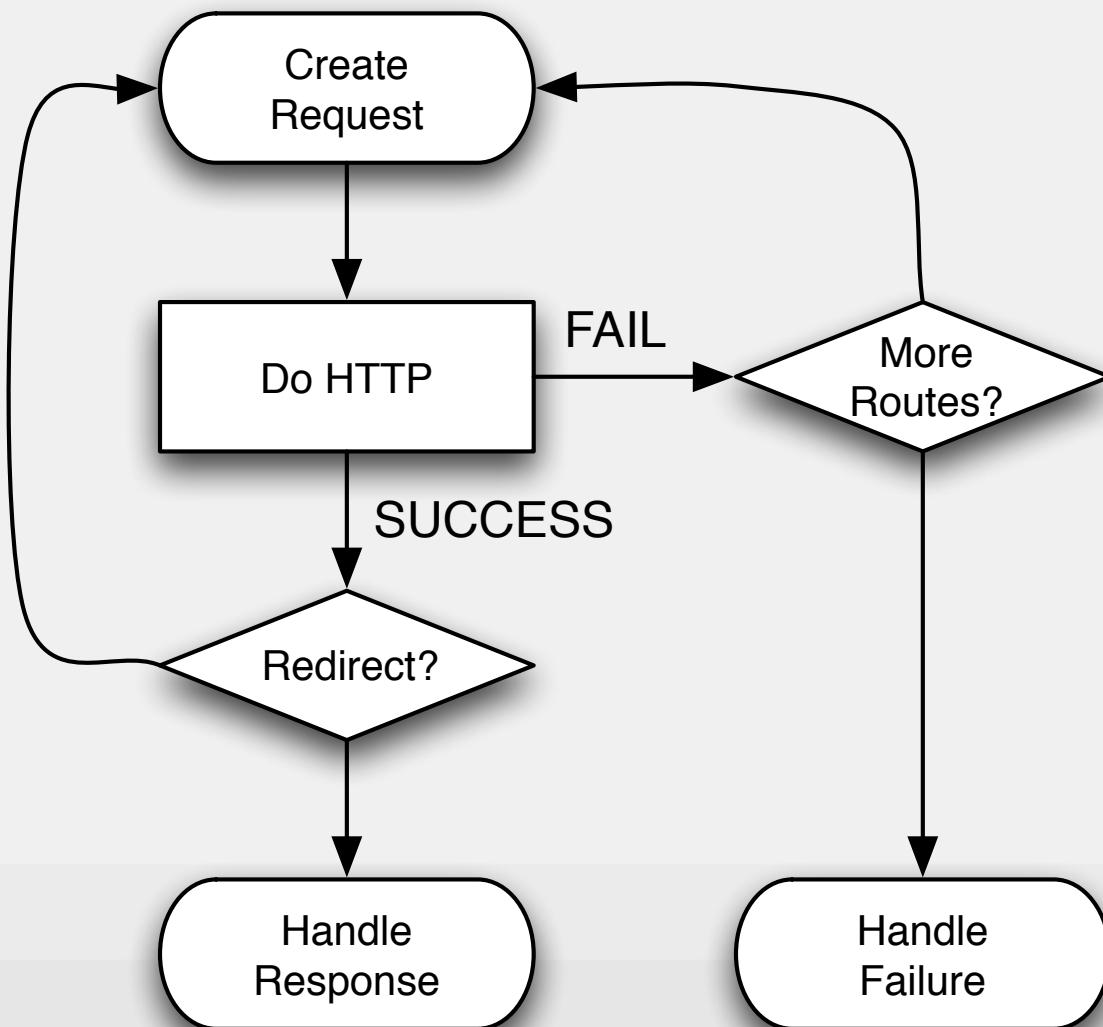
---



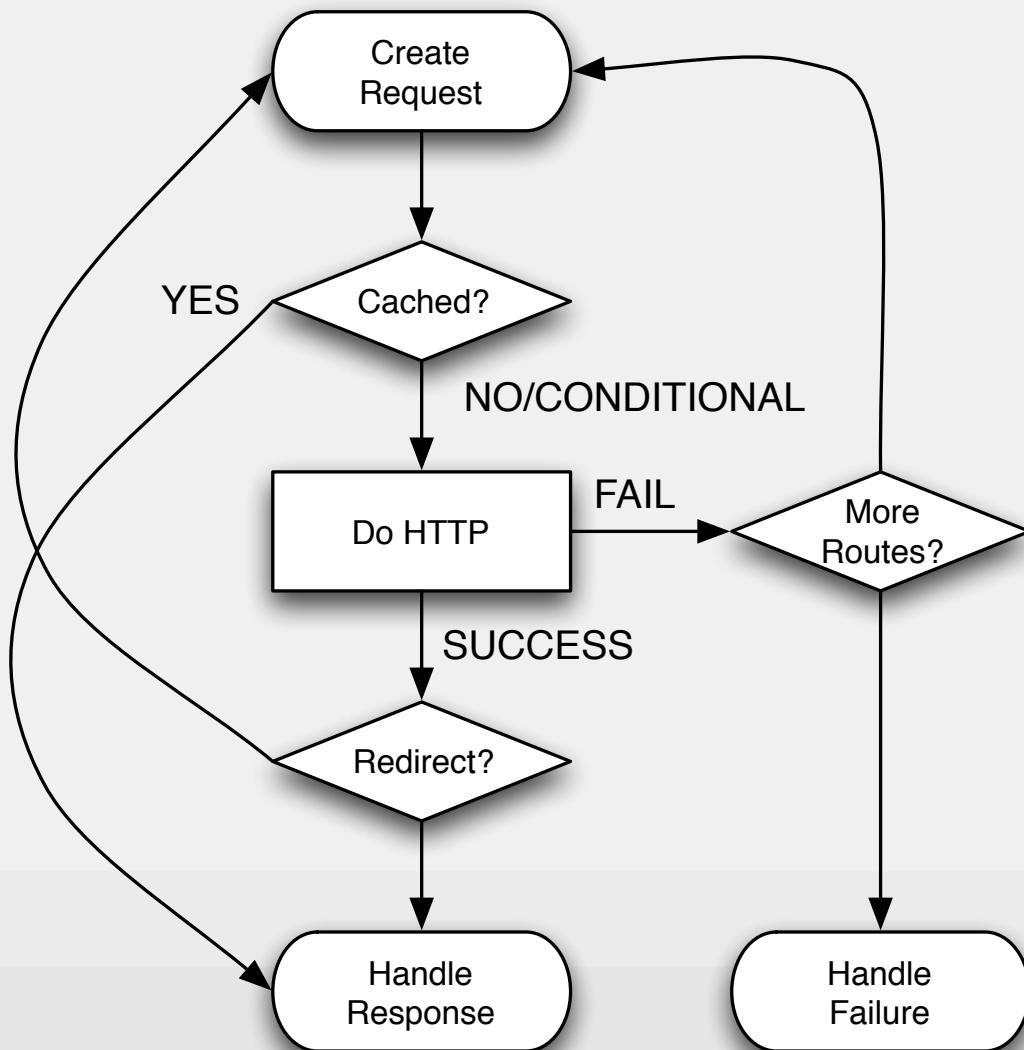
HTTP/1.1 302 Moved  
Location: /account

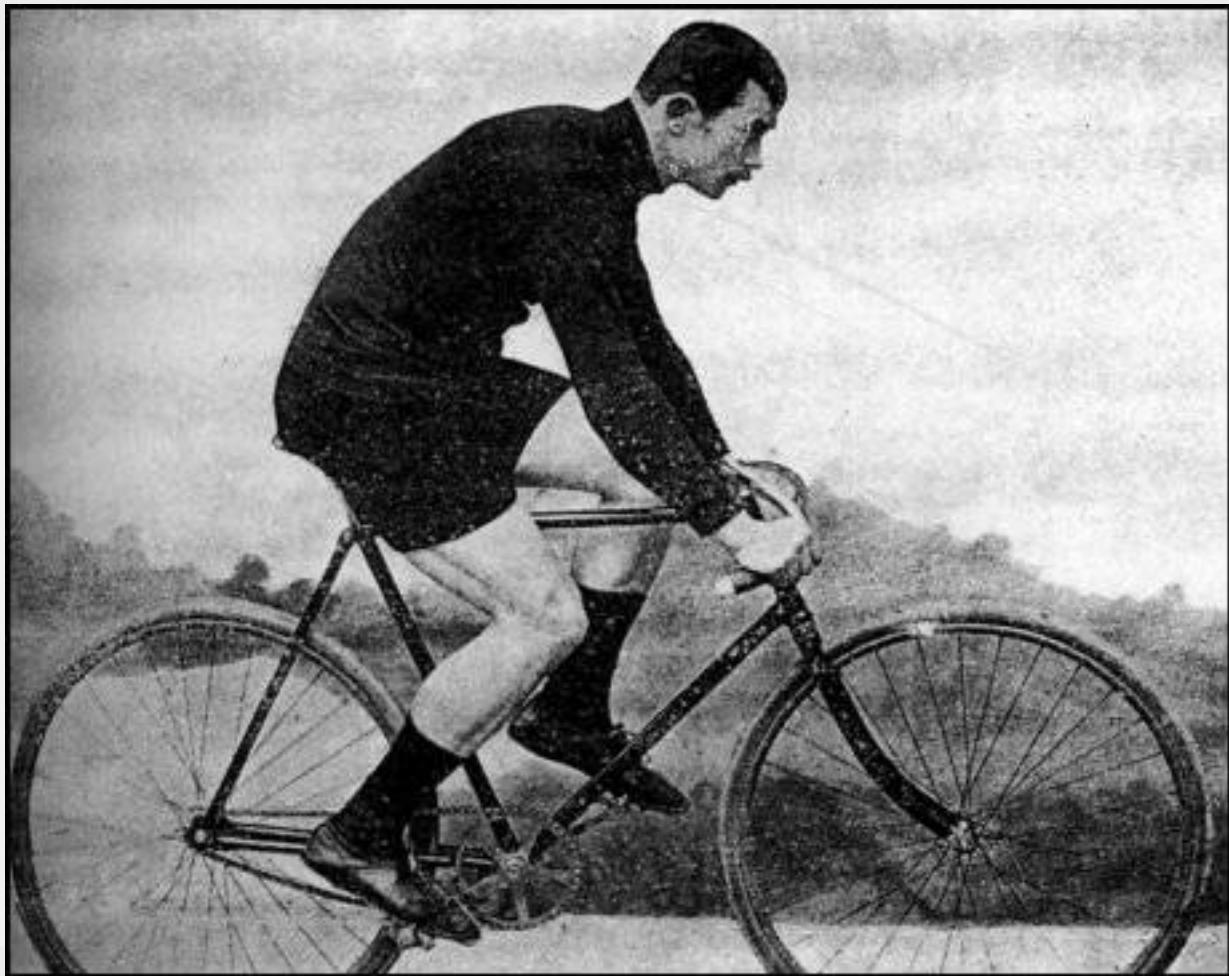
# HTTP with redirects and routing

---



# HTTP with redirects, routing and caching





<http://www.flickr.com/photos/paukrus/>

# HttpURLConnection

---

- Limited API, grown organically since the 1996
- ~ 9,000 lines of code

# HttpURLConnection

---

- Problems
  - Connection pool configuration is gross system properties
  - Connection pool is buggy in Froyo and should be disabled

# Demo



<http://www.flickr.com/photos/dhwright>

# Apache HTTP Client

---

- Mature & capable
- Lots of knobs & configuration
- ~ 50,000 lines of code

# Apache HTTP Client

---

- Problems:
  - Where does the work happen? How do I get the behavior I want?
  - Enterprisey: `HttpAbstractParamBean`

# Demo



<http://www.flickr.com/photos/turbulentflow>

# OkHttp

---

- Forked from Android 4.0's HttpURLConnection
- New features:
  - SPDY
  - Fault recovery
- ~ 12,000 lines of code



# OkHttp

---

- Problems:
  - It's new. No deep archive of support on stackoverflow.com!
  - Still uses the old API



# Demo

# Essential Dev Tools

# Charles

---

- Web Debugging Proxy
- \$50

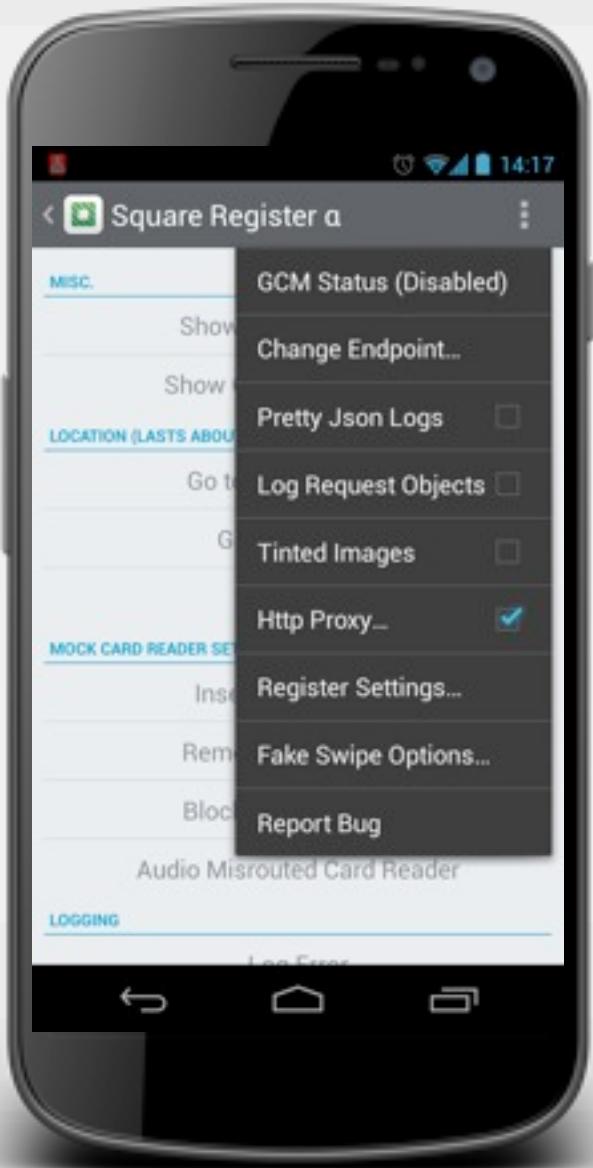


<http://www.charlesproxy.com>

# Demo

# Charles

---



# Charles

The screenshot shows the Charles 3.6.5 interface. The title bar reads "Charles 3.6.5 – Session 1 \*". The toolbar contains icons for file operations (New, Open, Save, Delete, Find, Copy, Paste, Find in Content, Stop, Reload, Stop All, Reload All), and various tools (Copy to Clipboard, Edit, Inspect, Verify, Tools, Settings).

The left pane displays a tree view of the request structure:

- https://api.broadway.squareup.com
  - 1.0/
    - items/
      - sync
      - sync
      - sync
      - sync
    - account/
      - status/
  - preferences

The "preferences" item is highlighted with a green selection bar.

The right pane shows the response details for the selected request:

HTTP/1.1 200 OK

Content-Type: application/json; charset=utf-8  
Transfer-Encoding: chunked  
Connection: keep-alive  
Keep-Alive: timeout=300  
Status: 200  
Cache-Control: private, max-age=0, must-revalidate  
ETag: "bc70ceba41a22a5e4c5ea6be5c9ba114"  
X-Frame-Options: DENY  
X-Square: S=stage02.mtv.squareup.com  
Strict-Transport-Security: max-age=1296000  
X-Content-Type-Options: nosniff  
X-XSS-Protection: 1; mode=block  
Content-Encoding: gzip

```
{"tax_enabled":false,"tipping_use_manual_tip_entry":false,"skip_si
```

# MockWebServer

---

- Does my HTTP code work?



# Mocks?

---

EASYMock

jMock



# Mocks!

---

It's a script!

1. Describe a scenario.
2. Exercise your code.
3. Verify.

# Demo

# Retrofit



# Retrofit

---

- Manually calling HTTP services requires a lot of code
- There's pain!
  - Error handling
  - Async
  - Testing

# Retrofit

---

- Adapts an HTTP service to a Java interface



# Demo

# Picasso



# Picasso

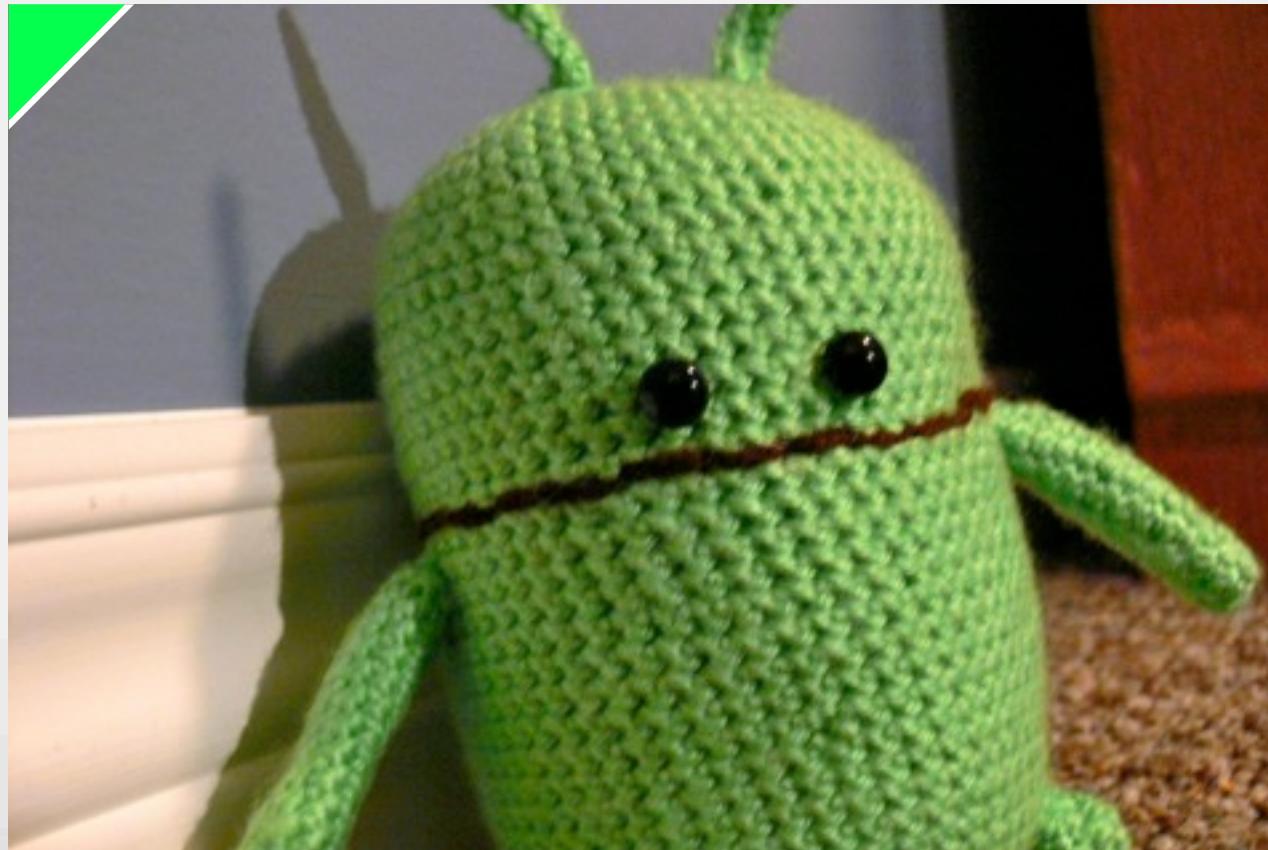
---

- Manually showing images on Android requires a lot of code
- There's pain!
  - Threading
  - View recycling
  - Memory caching
  - Disk caching
  - Transforming images

# Picasso

---

- Makes it foolishly easy to show images in your Android app



<http://www.flickr.com/photos/jumpyjodes>

# Demo

Volley

# Volley

---

- An HTTP library from the guys who make the Play Store app
- Sophisticated features
  - Request canceling
  - Request scheduler
  - Pluggable

# Demo

# Wrapping Up

## Advice: HTTP

---

- HTTP has knobs
- Use them

## Advice: HTTP

---

- Don't forget about the server!

# Advice: Development

---

- Don't fly blind!
  - Charles Web Debugging Proxy
  - MockWebServer

## Advice: Android

---

- High level libraries exist. Use them.
  - Retrofit, Picasso, Volley

# Questions?

[squareup.com/careers](http://squareup.com/careers)

